

Capstone Project - WALMART SALES PREDICTION PROJECT

-by Shivani Gupta

Table of Contents

1. Problem Statement
2. Project Objective
3. Data Description
4. Data Pre-processing Steps and Inspiration
5. Choosing the Algorithm for the Project
6. Motivation and Reasons For Choosing the Algorithm
7. Model Evaluation and Techniques
8. Inferences from the Same
9. Future Possibilities of the Project
10. References

Problem Statement

A retail store that has multiple outlets across the country are facing issues in managing the inventory - to match the demand with respect to supply. You are a data scientist, who has to come up with useful insights using the data and make prediction models to forecast the sales for X number of months/years.

Project Objective

1. Using the above data, come up with useful insights that can be used by each of the stores to improve in various areas.
2. Forecast the sales for each store for the next 12 weeks.

Data Description

The walmart.csv dataset contains 6435 rows and 8 columns. The range of the data Date column varies from Jan 2010 to Dec 2012.

Feature Name	Description
Store	Store number
Date	Week of Sales
Weekly_Sales	Sales for the given store in that week
Holiday_Flag	If it is a holiday week
Temperature	Temperature on the day of the sale
Fuel_Price	Cost of the fuel in the region
CPI	Consumer Price Index
Unemployment	Unemployment Rate

Data Preprocessing Steps And Inspiration

Checking for the insights into data:

```
In [70]: warnings.filterwarnings("ignore", category=FutureWarning)
walmart_data.describe(include='all')
```

Out[70]:

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment	Unemployment_treated	Year	Month
count	6434.000000	6434	6.434000e+03	6434.000000	6434.000000	6434.000000	6434.000000	6434.000000	6434.000000	6434.000000	6434
unique	NaN	143	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	12
top	NaN	2010-05-02 00:00:00	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	April
freq	NaN	45	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	629
first	NaN	2010-01-10 00:00:00	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
last	NaN	2012-12-10 00:00:00	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
mean	23.002487	NaN	1.047041e+06	0.069941	60.673531	3.358661	171.575257	7.999024	7.871061	2010.965030	NaN
std	12.987660	NaN	5.643776e+05	0.255067	18.429779	0.459035	39.358967	1.876003	1.520766	0.797081	NaN
min	1.000000	NaN	2.099862e+05	0.000000	5.540000	2.472000	126.064000	3.879000	4.294500	2010.000000	NaN
25%	12.000000	NaN	5.531677e+05	0.000000	47.495000	2.933000	131.735000	6.891000	6.891000	2010.000000	NaN
50%	23.000000	NaN	9.608457e+05	0.000000	62.675000	3.446500	182.616521	7.874000	7.874000	2011.000000	NaN
75%	34.000000	NaN	1.420282e+06	0.000000	74.945000	3.735000	212.745096	8.622000	8.622000	2012.000000	NaN
max	45.000000	NaN	3.818686e+06	1.000000	100.140000	4.468000	227.232807	14.313000	11.218500	2012.000000	NaN

- **Convert Data Types:** Ensure that the data types of variables are appropriate for their respective columns.

```
In [4]: walmart_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6435 entries, 0 to 6434
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Store           6435 non-null   int64
1   Date            6435 non-null   object
2   Weekly_Sales    6435 non-null   float64
3   Holiday_Flag    6435 non-null   int64
4   Temperature     6435 non-null   float64
5   Fuel_Price      6435 non-null   float64
6   CPI             6435 non-null   float64
7   Unemployment    6435 non-null   float64
dtypes: float64(5), int64(2), object(1)
memory usage: 402.3+ KB
```

Since Date column datatype is object so lets change it to DateTime.

```
In [71]: warnings.filterwarnings("ignore", category=UserWarning)
walmart_data['Date'] = walmart_data['Date'].astype('datetime64[ns]')
```

```
In [6]: walmart_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6435 entries, 0 to 6434
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   Store            6435 non-null   int64   
1   Date             6435 non-null   datetime64[ns]
2   Weekly_Sales     6435 non-null   float64  
3   Holiday_Flag     6435 non-null   int64   
4   Temperature      6435 non-null   float64  
5   Fuel_Price       6435 non-null   float64  
6   CPI              6435 non-null   float64  
7   Unemployment     6435 non-null   float64  
dtypes: datetime64[ns](1), float64(5), int64(2)
memory usage: 402.3 KB
```

- **Handle Missing Values:** Check for missing/Null values in the dataset and decide on the appropriate strategy for handling them.

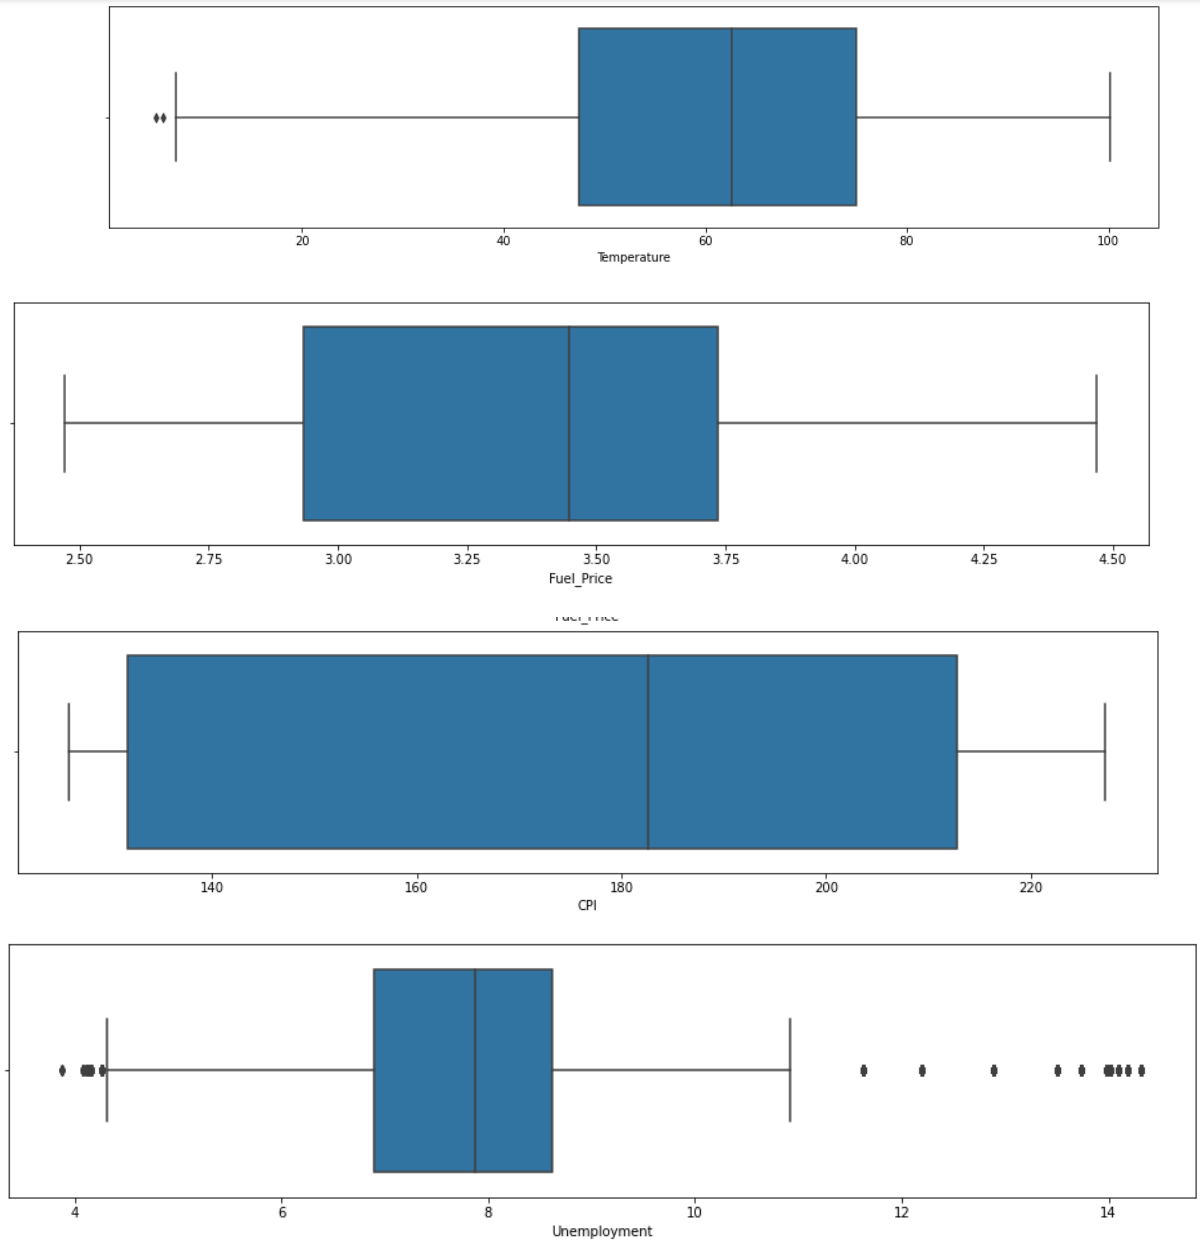
```
In [10]: walmart_data.isnull().sum()
```

```
Out[10]: Store            0
Date                    0
Weekly_Sales           0
Holiday_Flag           0
Temperature            0
Fuel_Price             0
CPI                   0
Unemployment           0
dtype: int64
```

- **Remove Outliers:** Identify outliers in the data that may adversely affect the analysis or model performance. Decide whether to remove outliers or apply appropriate transformations to mitigate their impact.

Checking for outliers

```
In [73]: fig, axis = plt.subplots(4, figsize=(16,16))
x=walmart_data[['Temperature', 'Fuel_Price', 'CPI', 'Unemployment']]
for i, column in enumerate(x):
    sns.boxplot(walmart_data[column], ax=axis[i])
warnings.filterwarnings("ignore", category=FutureWarning)
```



We can see that there are outliers for column Unemployment and Temperature. Checking the number of data rows that have outliers and deciding whether to delete them or not.


```
In [12]: walmart_data[(walmart_data['Unemployment']<4.5) | (walmart_data['Unemployment']>11)].shape
Out[12]: (520, 8)
```

```
# Since data is larger so we cannot drop this. Treating the outliers here.
```

Since data is larger so we cannot drop this. Treating the outliers here.

```
In [13]: # Using IQR Method
def thr_min_max(col):
    p25= walmart_data[col].quantile(0.25)
    p75=walmart_data[col].quantile(0.75)
    IQR=p75-p25

    thr_min,thr_max = p25 - 1.5* IQR , p75+1.5*IQR
    return thr_min,thr_max

def treating(val):
    if(val<thr_min):
        return thr_min
    elif(val>thr_max):
        return thr_max
    else:
        return val
```

```
In [14]: thr_min,thr_max = thr_min_max('Unemployment')
print(thr_min,thr_max)
walmart_data['Unemployment_treated'] = walmart_data['Unemployment'].apply(treating)

4.2945 11.218499999999999
```

So, Unemployment data is treated and new column is created for the same named Unemployment_treated.

```
In [78]: walmart_data.sample(7)
```

```
Out[78]:
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment	Unemployment_treated	Year	Month	Day
5360	38	2011-03-06	396826.06	0	71.81	4.031	129.049032	13.736	11.2185	2011	March	6
2260	16	2012-04-20	436221.26	0	43.61	3.936	197.722738	6.169	6.1690	2012	April	20
3703	26	2012-07-20	1049625.90	0	66.75	3.610	138.233193	7.405	7.4050	2012	July	20
3442	25	2010-04-16	715311.60	0	52.16	2.899	203.730749	7.856	7.8560	2010	April	16
5503	39	2011-03-06	1541745.59	0	83.59	3.699	214.016280	8.300	8.3000	2011	March	6
5526	39	2011-11-11	1456957.38	0	63.11	3.297	216.721737	7.716	7.7160	2011	November	11
2651	19	2011-07-29	1298775.80	0	74.86	4.004	135.963935	7.806	7.8060	2011	July	29

Here, we can see that in Unemployment column, value was larger than 11 so we had it replaced with threshold maximum value.

Now, checking number of outlier rows for Temperature column.

```
In [15]: walmart_data[walmart_data['Temperature']<5]
Out[15]:
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment	Unemployment_treated
910	7	2011-04-02	558027.77	0	-2.06	3.011	191.762589	8.818	8.818

```
In [16]: # Only one column with outlier of Temperature. So dropping it off.
walmart_data = walmart_data.drop(walmart_data[walmart_data['Temperature']<5].index)
```

Since there was only one row so instead of treating it we can delete it directly and it will not have any impact on our further observations.

- **Feature Engineering:** Explore the dataset to identify any potential feature engineering opportunities. This may involve creating new variables, extracting useful information from existing variables, or transforming variables to enhance their predictive power.

```
In [18]: walmart_data['Store'].unique()
# 45 stores data is present
Out[18]: array([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16, 17,
                18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34,
                35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45], dtype=int64)
```

Here, we can see that there are total of 45 stores data present in the dataset.

Checking which store has maximum sales:

```
In [79]: # Which store has maximum sales
```

```
px.bar(data_frame=walmart_data,x='Store',y='Weekly_Sales')
```




Here we can see store 20 has max sales and then Store 4 has 2nd highest sale.

Checking which store has maximum standard deviation:

```
# store having maximum standard deviation
```

```
store_std = walmart_data.groupby('Store')['Weekly_Sales'].std().reset_index()
```


```
px.bar(x=store_std['Store'],y=store_std['Weekly_Sales'])
```



Thus, the store which has maximum standard deviation is store number 14, which means that the sales of Store 14 varies the most and is not constant.

Checking which month has highest sales:


```
px.bar(x=walmart_data['Date'],y=walmart_data['Weekly_Sales'])
```



We can observe from the above graph that sales for the month of December are the maximum sales in a year. So, during the December months, year-end sales and other sales attract the customer more and also, the number of holidays at that time is more. So, the total sales in that month increase as compared to the rest of the year.

Checking which Store has maximum sales for Year 2010:

```
In [26]: data = walmart_data[(walmart_data['Date'] > '2010-01-01') & (walmart_data['Date'] < '2011-01-01')].groupby('Store')
[Weekly_Sales].sum()
data = pd.DataFrame(data).reset_index()
px.bar(x=data['Store'], y=data['Weekly_Sales'])
```



We can see that store 14 has maximum Sales in Year 2010.

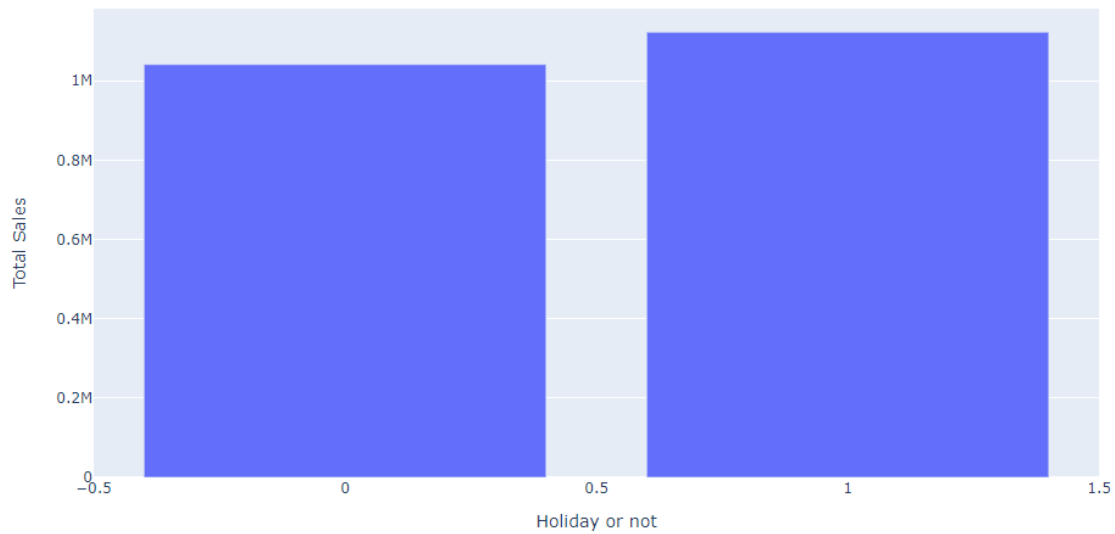
Check whether the Holidays have effect on Sales or not

```
In [27]: # Checking the effect of holiday on Sales
data = walmart_data.groupby('Holiday_Flag')['Weekly_Sales'].mean()
data = pd.DataFrame(data).reset_index()
data
```

```
Out[27]:
```

	Holiday_Flag	Weekly_Sales
0	0	1.041337e+06
1	1	1.122888e+06

```
In [82]: warnings.filterwarnings("ignore", category=DeprecationWarning)
px.bar(x=data['Holiday_Flag'], y=data['Weekly_Sales'], labels={'x': 'Holiday or not', 'y': 'Total Sales'})
```




Here, we can see that Sales is higher during holidays.

Checking Year Wise Monthly Sales:


In [30]: *#Year-wise Monthly Sales*

```
walmart_data['Year'] = pd.to_datetime(walmart_data['Date']).dt.year
walmart_data['Month'] = pd.to_datetime(walmart_data['Date']).dt.month_name()
walmart_data['Day'] = pd.to_datetime(walmart_data['Date']).dt.day
year_month_sales = walmart_data.groupby(['Year', 'Month'])['Weekly_Sales'].sum().reset_index()
print(year_month_sales.sort_values(by=['Year', 'Month']))
px.bar(data_frame=year_month_sales, x='Month', y='Weekly_Sales', color='Year')
```



Check effect of Fuel_Price on Weekly_Sales

```
px.scatter(x=walmart_data['Fuel_Price'],y=walmart_data['Weekly_Sales'])
```



From this we can say there is no direct relation between fuel price and sales.

- **Feature Scaling:** Normalize or standardize numerical features.

```
In [38]: # Standardization
```

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X = walmart_data[['Store', 'Fuel_Price', 'CPI', 'Unemployment', 'Year']]
Y = walmart_data['Weekly_Sales']
X_std = sc.fit_transform(X)
X_std
```

```
Out[38]: array([[ -1.69423879,  -1.7138609 ,  1.00419743,  0.05702792, -1.21079834],
 [ -1.69423879,  -1.76614857,  1.00790238,  0.05702792, -1.21079834],
 [ -1.69423879,  -1.84022277,  1.00909593,  0.05702792, -1.21079834],
 ...,
 [ 1.69385582,  1.39725531,  0.52729082,  0.35609123,  1.29855149],
 [ 1.69385582,  1.32971707,  0.527382 ,  0.35609123,  1.29855149],
 [ 1.69385582,  1.14017428,  0.52682414,  0.35609123,  1.29855149]])
```

Choosing the Algorithm for Walmart Data

I have chosen Random Forest Regressor Algorithm for this project because after checking multiple algorithms we have maximum Accuracy and least error in Random Forest.

I have chosen the Random Forest Regressor algorithm for this project because, after checking multiple algorithms, it has shown the highest accuracy and the least error compared to other algorithms.

Here are some reasons why I have chosen the Random Forest Regressor algorithm:

- **High Accuracy:** In my evaluation of various algorithms, the Random Forest Regressor has demonstrated the highest accuracy in predicting the target variable. It has shown better performance in capturing complex relationships between the features and the target, resulting in more accurate predictions.
- **Robustness:** Random Forest Regressor is robust against overfitting and tends to generalize well to unseen data. By constructing multiple decision trees and averaging their predictions, it reduces the impact of individual noisy or biased trees, leading to improved robustness and more reliable predictions.
- **Feature Importance:** Random Forest Regressor provides valuable insights into feature importance. It calculates the importance score of each feature, indicating its contribution to the prediction. This information can help identify the most influential features in the dataset, allowing for better understanding and interpretation of the underlying patterns.
- **Handling Nonlinear Relationships:** The Random Forest Regressor algorithm can effectively capture nonlinear relationships between the features and the target variable. It does not assume linearity and can handle complex interactions and dependencies among features, making it suitable for datasets with intricate patterns.
- **Robustness to Outliers and Missing Values:** Random Forest Regressor is robust to outliers and can handle missing values without requiring imputation. It uses the information available in other features to make predictions, reducing the impact of missing or outlier values on the overall model performance.

By selecting the Random Forest Regressor algorithm, I aim to leverage its strengths to build a reliable and accurate prediction model for the given dataset. It has demonstrated superior performance in terms of accuracy and error reduction, making it a suitable choice for this project.

LINEAR REGRESSION MODEL:


```
In [39]: # Linear Regression :  
  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression  
X_train, X_test, Y_train, Y_test = train_test_split(X_std, Y, test_size=0.2)  
X_train.shape, X_test.shape, Y_train.shape, Y_test.shape
```

```
Out[39]: ((5147, 5), (1287, 5), (5147,), (1287,))
```

```
In [41]: reg = LinearRegression()  
reg.fit(X_train, Y_train)  
Y_test_pred = reg.predict(X_test)  
from sklearn import metrics  
print('Accuracy of test data:', reg.score(X_test, Y_test)*100)  
print('Mean Absolute Error:', metrics.mean_absolute_error(Y_test, Y_test_pred))  
print('Mean Squared Error:', metrics.mean_squared_error(Y_test, Y_test_pred))  
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(Y_test, Y_test_pred)))  
sns.scatterplot(Y_test_pred, Y_test)
```

```
Accuracy of test data: 16.872680299024058  
Mean Absolute Error: 430045.12184525817  
Mean Squared Error: 270162159436.3094  
Root Mean Squared Error: 519771.25683930365
```

```
Out[41]: <AxesSubplot:ylabel='Weekly_Sales'>
```




The accuracy is very less. So we cannot say this is a good model.

DECISION TREE MODEL:

```
In [43]: # Decision Tree
from sklearn.tree import DecisionTreeRegressor
X_train, X_test, y_train, y_test = train_test_split(X_std, Y, test_size=0.2)
model = DecisionTreeRegressor()
model.fit(X_train, y_train)
y_pred_train = model.predict(X_train)
y_pred_test = model.predict(X_test)
print("Accuracy of Train data:", model.score(X_train, y_train)*100)
print("Accuracy of test data:", model.score(X_test, y_test)*100)
print("Mean Absolute Error: ", metrics.mean_absolute_error(y_test, y_pred_test))
print("Mean Square Error: ", metrics.mean_squared_error(y_pred_test, y_test))
print("Root Mean Square Error: ", np.sqrt(metrics.mean_squared_error(y_pred_test, y_test)))
sns.scatterplot(y_test, y_pred_test)
```

```
Accuracy of Train data: 100.0
Accuracy of test data: 88.32363820734517
Mean Absolute Error: 92774.59034965035
Mean Square Error: 36454446134.767
Root Mean Square Error: 190930.47460991394
```

Out[43]: <AxesSubplot:xlabel='Weekly_Sales'>




RANDOM FOREST MODEL:

```
In [89]: # Random Forest Regressor
from sklearn.ensemble import RandomForestRegressor
randomfrst = RandomForestRegressor()
randomfrst.fit(X_train, Y_train)
Y_pred = randomfrst.predict(X_test)
print('Accuracy of test data:', randomfrst.score(X_test, Y_test)*100)
print('Mean Absolute Error:', metrics.mean_absolute_error(Y_test, Y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(Y_test, Y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(Y_test, Y_pred)))
sns.scatterplot(Y_pred, Y_test)
```

```
Accuracy of test data: 93.65615561560301
Mean Absolute Error: 72084.45893123545
Mean Squared Error: 19065726429.965282
Root Mean Squared Error: 138078.69651023392
```

```
Out[89]: <AxesSubplot:ylabel='Weekly_Sales'>
```



Here, Accuracy of Linear regression = 10%.

Decision Tree Model = 87%

Random Forest Model = 93%

So, Accuracy and error of Random Forest Regression Model is highest and best, So we will be choosing this model.

Model Evaluation and Technique

Model evaluation is a crucial step in machine learning to assess the performance and effectiveness of a predictive model. It involves using various techniques and metrics to measure how well the model generalizes to new, unseen data and how accurately it predicts the target variable. Here are model evaluation techniques and metrics used:

★ **Train-Test Split:** This technique involves splitting the available data into two subsets: a training set and a test set. The model is trained on the training set and then evaluated on the test set to assess its performance. The accuracy and error metrics on the test set provide an indication of how well the model is expected to perform on new, unseen data.

★ **Metrics for Regression Models:**

- Mean Squared Error (MSE): Measures the average squared difference between the predicted and actual values. A lower MSE indicates better performance. We got the lowest MSE for Random Forest Model.
- Root Mean Squared Error (RMSE): The square root of MSE, which provides a measure of the average prediction error in the original units of the target variable. We got the lowest RMSE for Random Forest Model.
- Mean Absolute Error (MAE): Measures the average absolute difference between the predicted and actual values. It is less sensitive to outliers compared to MSE. We got the lowest MAE for Random Forest Model.

★ **Metrics for Classification Models:**

- Accuracy: Measures the proportion of correctly classified instances. It is the most commonly used metric for classification tasks.

```
Accuracy of test data: 94.06832847178373
Mean Absolute Error: 74680.64108453768
Mean Squared Error: 19567181988.359875
Root Mean Squared Error: 139882.74371186702
```

★ **Predicting Values for future Sale per Store:**

To predict Future Sales I have used Time Series model. Checking whether the data is stationary or not.


```
In [54]: from statsmodels.tsa.stattools import adfuller
sales_data = walmart_data['Weekly_Sales']
result = adfuller(sales_data)
p_value = result[1]
print("P_value: ",p_value)
if p_value < 0.05:
    print("The data is stationary.")
else:
    print("The data is not stationary.")
```

```
P_value: 0.00011613258802178222
The data is stationary.
```

To check if the data is Seasonal or not:

```
In [51]: walmart_data['Weekly_Sales'].plot()|
# This is seasonal data
```

```
Out[51]: <AxesSubplot:>
```



Here since the plot is symmetric and repeating itself so we can conclude that our data is seasonal. Since our data is symmetric and stationary then we can directly use SARIMA Time Series model to predict the sales for next 12 months.

```
In [62]: import warnings

with warnings.catch_warnings():
    warnings.filterwarnings("ignore", category=DeprecationWarning)

import statsmodels.api as sm

unique_stores = walmart_data['Store'].unique()

order = (1, 1, 1)
seasonal_order = (1, 1, 1, 12)

for store_id in unique_stores:
    store_data = walmart_data[walmart_data['Store'] == store_id]
    model = sm.tsa.statespace.SARIMAX(store_data['Weekly_Sales'], order=order, seasonal_order=seasonal_order)
    model_fit = model.fit()
    forecast = model_fit.forecast(steps=12)
    print(f"Store ID: {store_id}")
    print(forecast)
    print("-----")
```

For the above code, we got the forecast store wise for next 12 months.

Store ID: 1		Store ID: 2	
143	1.421157e+06	143	1.709647e+06
144	1.511237e+06	144	1.818715e+06
145	1.498984e+06	145	1.833277e+06
146	1.571846e+06	146	1.946534e+06
147	1.510895e+06	147	1.790230e+06
148	1.472766e+06	148	1.757251e+06
149	1.497334e+06	149	1.805713e+06
150	1.504977e+06	150	1.806947e+06
151	1.498639e+06	151	1.796017e+06
152	1.497632e+06	152	1.812438e+06
153	1.511394e+06	153	1.827277e+06
154	1.588019e+06	154	1.938755e+06
Name: predicted_mean, dtype: float64		Name: predicted_mean, dtype: float64	
-----		-----	

Store ID: 3

143 388216.857328
144 404162.782058
145 412566.886013
146 413998.263833
147 404024.811650
148 397529.323935
149 410445.089803
150 399799.721293
151 415801.653782
152 408450.567700
153 416225.871699
154 425200.567853

Name: predicted_mean, dtype: float64

Store ID: 5

143 301629.058840
144 316485.721285
145 322151.589739
146 332790.165539
147 324724.096373
148 305834.757204
149 314847.238615
150 322344.244309
151 322412.734600
152 319636.809249
153 319885.382386
154 337477.321497

Name: predicted_mean, dtype: float64

Store ID: 7

142 483026.421967
143 476887.523339
144 505055.258846
145 558702.924950
146 524734.147861
147 499909.135404
148 499370.220851
149 502102.623655
150 474289.182954
151 468152.625244
152 472499.611932
153 517706.128679

Name: predicted_mean, dtype: float64

Store ID: 4

143 2.041562e+06
144 2.144343e+06
145 2.151268e+06
146 2.221072e+06
147 2.092061e+06
148 2.075226e+06
149 2.117922e+06
150 2.108841e+06
151 2.102557e+06
152 2.120777e+06
153 2.127436e+06
154 2.254843e+06

Name: predicted_mean, dtype: float64

Store ID: 6

143 1.326060e+06
144 1.393026e+06
145 1.420462e+06
146 1.514057e+06
147 1.404896e+06
148 1.321009e+06
149 1.373889e+06
150 1.398986e+06
151 1.343666e+06
152 1.362866e+06
153 1.382155e+06
154 1.474965e+06

Name: predicted_mean, dtype: float64

Store ID: 8

143 856652.166968
144 888352.748699
145 914389.420730
146 953022.732729
147 885269.091547
148 872350.628913
149 883006.468432
150 898874.555100
151 888383.687663
152 897320.218744
153 904271.287712
154 948172.661624

Name: predicted_mean, dtype: float64

Store ID: 9	Store ID: 10
143 517610.484030	143 1.589181e+06
144 538699.155897	144 1.684137e+06
145 560763.164998	145 1.684609e+06
146 577024.023170	146 1.746737e+06
147 541794.325176	147 1.630742e+06
148 532093.759780	148 1.583527e+06
149 542536.372482	149 1.626618e+06
150 545777.050051	150 1.662555e+06
151 550092.456763	151 1.620476e+06
152 546778.369030	152 1.668422e+06
153 548325.312332	153 1.696395e+06
154 585263.154930	154 1.820696e+06
Name: predicted_mean, dtype: float64	Name: predicted_mean, dtype: float64
-----	-----
Store ID: 11	Store ID: 43
143 1.141184e+06	143 605823.526953
144 1.228769e+06	144 642219.939020
145 1.263735e+06	145 630465.179414
146 1.309240e+06	146 632010.146259
147 1.206687e+06	147 643462.345616
148 1.193993e+06	148 636426.551157
149 1.228307e+06	149 625476.926051
150 1.235919e+06	150 627410.091171
151 1.223171e+06	151 641311.989911
152 1.220821e+06	152 631060.586234
153 1.221971e+06	153 635819.639759
154 1.279151e+06	154 635610.194316
Name: predicted_mean, dtype: float64	Name: predicted_mean, dtype: float64
-----	-----
Store ID: 44	Store ID: 45
143 343302.480489	143 688334.902057
144 357312.901781	144 722057.332235
145 362221.014641	145 736361.133701
146 366679.858277	146 782560.025048
147 357964.403436	147 727876.662515
148 367574.672280	148 702308.183734
149 362833.176926	149 713208.502955
150 373710.585780	150 735743.965671
151 362136.162249	151 696949.241392
152 365218.358786	152 711343.324672
153 358537.578101	153 727171.451288
154 379721.270066	154 786707.313401
Name: predicted_mean, dtype: float64	Name: predicted_mean, dtype: float64
-----	-----

Similarly we got the output for all 45 stores. So, we have successfully predicted the futures sales for next 12 months for each store respectively.

Inferences

Based on the Walmart project, here are some possible inferences and insights that can be drawn from the analysis:

- **Seasonality:** The sales data from Walmart exhibits a clear pattern of seasonality, with higher sales during specific periods of the year. This can be further analyzed to identify the peak seasons and plan inventory, promotions, and staffing accordingly.
- **Store Performance:** The analysis of individual store sales data can provide insights into the performance of each store. Identifying stores with consistently high or low sales can help in understanding the factors contributing to their performance and implementing strategies to improve sales.
- **Promotions and Discounts:** Analyzing the impact of promotions and discounts on sales can help identify effective marketing strategies. By studying the correlation between promotional activities and sales spikes, Walmart can optimize its promotional campaigns and maximize their impact on sales.
- **Pricing Optimization:** Analyzing the relationship between pricing and sales can help identify optimal price points for different products or categories. This analysis can assist in pricing strategies, discount planning, and competitive positioning.
- **Store Layout and Placement:** Analyzing sales data along with store layout and product placement information can reveal insights into the impact of store design on customer behavior. It can help optimize store layouts to enhance customer experience and maximize sales.
- **Forecasting:** By applying time series analysis techniques and predictive models to the sales data, Walmart can forecast future sales and demand. This information can support inventory planning, production scheduling, and supply chain management.

Future Possibilities

The future possibilities for Walmart sales prediction are vast, considering the advancements in data analytics, machine learning, and technology. Here are some potential future possibilities for Walmart sales prediction:

- **Advanced Forecasting Models:** Develop more sophisticated forecasting models using advanced techniques such as deep learning, ensemble methods, or hybrid models that combine multiple algorithms. These models can provide more accurate and granular sales predictions at different levels, including individual stores, product categories, and geographical regions.
- **Real-Time Sales Prediction:** Implement real-time sales prediction systems that continuously analyze incoming data and provide up-to-date sales forecasts. This would enable Walmart to make timely decisions, such as inventory management, staffing adjustments, and pricing strategies, based on the most current sales predictions.
- **Integration of External Data:** Incorporate external data sources, such as social media trends, economic indicators, competitor data, or weather patterns, into the sales prediction models. By analyzing the impact of these external factors on sales, Walmart can gain deeper insights and make more informed business decisions.
- **Demand Sensing:** Use advanced demand sensing techniques to capture and analyze real-time customer demand signals from various sources, including point-of-sale data, online sales, customer reviews, and social media. By leveraging demand sensing capabilities, Walmart can respond quickly to changes in customer preferences, optimize inventory levels, and enhance overall supply chain efficiency.
- **Personalized Sales Prediction:** Develop personalized sales prediction models that consider individual customer behavior, preferences, and purchasing history. By tailoring sales predictions to each customer, Walmart can provide personalized offers, recommendations, and promotions, leading to increased customer satisfaction and loyalty.

- **Integration with IoT and Sensor Data:** Leverage Internet of Things (IoT) devices and sensor data within stores to capture and analyze real-time data on foot traffic, product placements, and customer interactions. This data can be used to enhance sales prediction models, optimize store layouts, improve product placements, and enhance the overall in-store customer experience.
- **Geo-Spatial Analysis:** Apply geo-spatial analysis techniques to understand regional variations in sales patterns. By incorporating geographical factors such as population density, demographics, and local market dynamics into sales prediction models, Walmart can better target specific regions with tailored marketing strategies and optimize product assortments based on regional preferences.
- **Integration with E-commerce:** Integrate sales prediction models with e-commerce platforms to provide accurate sales forecasts for online sales channels. This would enable Walmart to optimize online inventory, improve order fulfillment, and enhance the overall online shopping experience for customers.

These are just a few potential future possibilities for Walmart sales prediction. As technology continues to evolve and new data analysis techniques emerge, Walmart can leverage these advancements to further enhance its sales prediction capabilities and stay at the forefront of retail analytics.

References

1. <https://chat.openai.com/>
2. <https://www.walmart.com/>
3. <https://www.kaggle.com/>
4. <https://towardsdatascience.com>
5. <https://www.analyticsvidhya.com>
6. <https://github.com/>
7. <https://scikit-learn.org/>
8. <https://plotly.com/>
9. <https://matplotlib.org/>